

Experiment

November 29, 2025

1 Appendix A

```
[1]: import pandas as pd  
import pycaret  
import numpy as np  
import copy
```

1.1 Read Raw Data

```
[2]: df_raw = pd.read_csv('training_data_ht2025.csv')
```

```
[3]: df_raw
```

```
[3]:      hour_of_day  day_of_week  month  holiday  weekday  summertime  temp  \  
0              5          5     1        0        0            0    -7.2  
1             21          4     1        0        1            0    -1.3  
2             21          3     8        0        1            1   26.9  
3              1          6     1        0        0            0    3.1  
4             17          0     3        0        1            0   11.7  
...           ...        ...    ...      ...      ...        ...  ...  
1595            3          5     6        0        0            0    21.5  
1596            14          0     6        0        1            1   23.2  
1597            13          0     3        0        1            1   13.9  
1598            14          5     3        0        0            1   11.7  
1599            22          6     2        0        0            0    4.2  
  
      dew  humidity  precip  snow  snowdepth  windspeed  cloudcover  \  
0   -15.0      53.68    0.000    0       0.0      16.3      31.6  
1   -12.8      40.97    0.000    0       0.0      23.9      85.7  
2    21.8      73.39    0.000    0       0.0       0.0      81.1  
3    -4.0      59.74    0.000    0       0.0      19.2       0.0  
4   -11.4      18.71    0.000    0       0.0      10.5      44.6  
...           ...        ...    ...      ...      ...        ...  ...  
1595   19.4      87.68    0.000    0       0.0      10.6      24.4  
1596   20.1      82.43    2.217    0       0.0       9.8      92.1  
1597   -2.2      32.93    0.000    0       2.0      18.2      79.3  
1598   -9.3      22.09    0.000    0       0.0       5.8      24.4
```

```
1599    1.7      84.11   1.081     0      0.0      21.9      97.4  
  
       visibility  increase_stock  
0            16.0  low_bike_demand  
1            16.0  low_bike_demand  
2            16.0  low_bike_demand  
3            16.0  low_bike_demand  
4            16.0  low_bike_demand  
...          ...  
1595        16.0  low_bike_demand  
1596        10.4  low_bike_demand  
1597        16.0  low_bike_demand  
1598        16.0  high_bike_demand  
1599        16.0  low_bike_demand  
  
[1600 rows x 16 columns]
```

1.1.1 Check Data Contents

```
[4]: df_raw['increase_stock'].value_counts()
```

```
[4]: increase_stock  
low_bike_demand      1312  
high_bike_demand     288  
Name: count, dtype: int64
```

Based on above code that shows the values on the target variable *increase_stock*, there's a problem of Imbalance Dataset, where the records with low bike demand are much more compared to records with high bike demands.

```
[5]: df_raw['holiday'].value_counts()
```

```
[5]: holiday  
0      1547  
1       53  
Name: count, dtype: int64
```

```
[6]: df_raw['weekday'].value_counts()
```

```
[6]: weekday  
1      1136  
0       464  
Name: count, dtype: int64
```

```
[7]: df_raw['summertime'].value_counts()
```

```
[7]: summertime  
1      1030
```

```
0      570  
Name: count, dtype: int64
```

```
[8]: df_raw['increase_stock'].value_counts()
```

```
[8]: increase_stock  
low_bike_demand    1312  
high_bike_demand   288  
Name: count, dtype: int64
```

```
[9]: df_raw['hour_of_day'].value_counts()
```

```
[9]: hour_of_day  
0      84  
16     83  
13     78  
7      77  
23     77  
17     72  
21     71  
14     70  
3      69  
4      69  
5      67  
1      67  
8      66  
12     66  
15     64  
10     63  
9      63  
6      62  
11     58  
19     57  
22     56  
20     55  
18     53  
2      53  
Name: count, dtype: int64
```

```
[10]: df_raw['day_of_week'].value_counts()
```

```
[10]: day_of_week  
4      242  
0      240  
5      233  
6      231  
2      223
```

```
3      220
1      211
Name: count, dtype: int64
```

```
[11]: df_raw['month'].value_counts()
```

```
[11]: month
3      147
11     141
4      140
6      137
7      136
12     136
5      133
9      131
2      131
1      128
8      121
10     119
Name: count, dtype: int64
```

```
[12]: df_raw['snow'].value_counts()
```

```
[12]: snow
0      1600
Name: count, dtype: int64
```

Note: The Snow Features only have one value: 0, and therefore isn't really useful for this dataset.
And In Feature selection we can drop them.

```
[13]: df_raw.isnull().sum()
```

```
[13]: hour_of_day      0
day_of_week       0
month            0
holiday          0
weekday          0
summertime       0
temp              0
dew               0
humidity          0
precip            0
snow              0
snowdepth         0
windspeed         0
cloudcover        0
visibility        0
increase_stock   0
```

```
dtype: int64
```

Based on the checking results, there're no NaN values, so we can proceed with feature processing.

Cyclical Encoding Function to encode Ordinal Variables into Cyclical encoding using Sine and Cosine

```
[14]: def cyclical_encode(x, max_val, start_val=0):
    x_arr = np.asarray(x, dtype=float)

    # If values are 1..period (e.g. months 1..12), shift to 0..period-1
    if not start_val:
        x_arr = x_arr - 1

    angle = 2 * np.pi * x_arr / max_val
    sin_x = np.sin(angle)
    cos_x = np.cos(angle)

    if np.isscalar(x):
        return float(sin_x), float(cos_x)

    return sin_x, cos_x
```

```
[ ]:
```

```
[15]: # Transform Cyclical Values for hour_of_day

# Transform Cyclical Values for hour_of_day
# Hour of day ranges from 0 to 23
df_hour_feat = pd.DataFrame(df_raw['hour_of_day'].apply(lambda x:cyclical_encode(x, max(df_raw['hour_of_day']), min(df_raw['hour_of_day'])))).to_list()
df_hour_feat.columns = ['hour_of_day_sin', 'hour_of_day_cos']
df_hour_feat

# Transform Cyclical Values for day_of_week
# Day of week ranges from 0 (Monday) to 6 (Sunday)
df_day_feat = pd.DataFrame(df_raw['day_of_week'].apply(lambda x:cyclical_encode(x, 6, 0))).to_list()
df_day_feat.columns = ['day_of_week_sin', 'day_of_week_cos']
df_day_feat

# Transform Cyclical Values for month
# Month ranges from 1 to 12
df_month_feat = pd.DataFrame(df_raw['month'].apply(lambda x: cyclical_encode(x, 12, 1))).to_list()
```

```

df_month_feat.columns = ['month_sin', 'month_cos']
df_month_feat

# Concat Cyclical Features
df_cyclical = pd.concat([df_hour_feat, df_day_feat, df_month_feat], axis=1)
df_cyclical

```

[15]:

	hour_of_day_sin	hour_of_day_cos	day_of_week_sin	day_of_week_cos	\
0	0.887885	0.460065	-8.660254e-01	-0.5	
1	-0.730836	0.682553	1.224647e-16	-1.0	
2	-0.730836	0.682553	8.660254e-01	-0.5	
3	0.000000	1.000000	-8.660254e-01	0.5	
4	-0.942261	-0.334880	-8.660254e-01	0.5	
...
1595	0.519584	0.854419	-8.660254e-01	-0.5	
1596	-0.398401	-0.917211	-8.660254e-01	0.5	
1597	-0.136167	-0.990686	-8.660254e-01	0.5	
1598	-0.398401	-0.917211	-8.660254e-01	-0.5	
1599	-0.519584	0.854419	-8.660254e-01	0.5	

	month_sin	month_cos
0	5.000000e-01	8.660254e-01
1	5.000000e-01	8.660254e-01
2	-8.660254e-01	-5.000000e-01
3	5.000000e-01	8.660254e-01
4	1.000000e+00	6.123234e-17
...
1595	1.224647e-16	-1.000000e+00
1596	1.224647e-16	-1.000000e+00
1597	1.000000e+00	6.123234e-17
1598	1.000000e+00	6.123234e-17
1599	8.660254e-01	5.000000e-01

[1600 rows x 6 columns]

```

[16]: # Creating Final Feature and Target DataFrames
df_features = copy.deepcopy(df_raw)

# Creating Target Variable and Mapping text to binary
df_target = df_features['increase_stock'].map({'low_bike_demand': 0, □
    ↵'high_bike_demand': 1})

# Dropping original cyclical columns and target from features
df_features = df_features.drop(['hour_of_day', 'day_of_week', 'month', □
    ↵'increase_stock'], axis=1)

# Concatenating Cyclical Features to Features DataFrame

```

```

df_features = pd.concat([df_cyclical, df_features], axis=1)

df_features

```

[16]:

	hour_of_day_sin	hour_of_day_cos	day_of_week_sin	day_of_week_cos	\			
0	0.887885	0.460065	-8.660254e-01	-0.5				
1	-0.730836	0.682553	1.224647e-16	-1.0				
2	-0.730836	0.682553	8.660254e-01	-0.5				
3	0.000000	1.000000	-8.660254e-01	0.5				
4	-0.942261	-0.334880	-8.660254e-01	0.5				
...				
1595	0.519584	0.854419	-8.660254e-01	-0.5				
1596	-0.398401	-0.917211	-8.660254e-01	0.5				
1597	-0.136167	-0.990686	-8.660254e-01	0.5				
1598	-0.398401	-0.917211	-8.660254e-01	-0.5				
1599	-0.519584	0.854419	-8.660254e-01	0.5				
	month_sin	month_cos	holiday	weekday	summertime	temp	dew	\
0	5.000000e-01	8.660254e-01	0	0	0	-7.2	-15.0	
1	5.000000e-01	8.660254e-01	0	1	0	-1.3	-12.8	
2	-8.660254e-01	-5.000000e-01	0	1	1	26.9	21.8	
3	5.000000e-01	8.660254e-01	0	0	0	3.1	-4.0	
4	1.000000e+00	6.123234e-17	0	1	0	11.7	-11.4	
...	
1595	1.224647e-16	-1.000000e+00	0	0	1	21.5	19.4	
1596	1.224647e-16	-1.000000e+00	0	1	1	23.2	20.1	
1597	1.000000e+00	6.123234e-17	0	1	1	13.9	-2.2	
1598	1.000000e+00	6.123234e-17	0	0	1	11.7	-9.3	
1599	8.660254e-01	5.000000e-01	0	0	0	4.2	1.7	
	humidity	precip	snow	snowdepth	windspeed	cloudcover	visibility	
0	53.68	0.000	0	0.0	16.3	31.6	16.0	
1	40.97	0.000	0	0.0	23.9	85.7	16.0	
2	73.39	0.000	0	0.0	0.0	81.1	16.0	
3	59.74	0.000	0	0.0	19.2	0.0	16.0	
4	18.71	0.000	0	0.0	10.5	44.6	16.0	
...	
1595	87.68	0.000	0	0.0	10.6	24.4	16.0	
1596	82.43	2.217	0	0.0	9.8	92.1	10.4	
1597	32.93	0.000	0	2.0	18.2	79.3	16.0	
1598	22.09	0.000	0	0.0	5.8	24.4	16.0	
1599	84.11	1.081	0	0.0	21.9	97.4	16.0	

[1600 rows x 18 columns]

[17]: df_target

```
[17]: 0      0
1      0
2      0
3      0
4      0
..
1595   0
1596   0
1597   0
1598   1
1599   0
Name: increase_stock, Length: 1600, dtype: int64
```

```
[18]: # Correlation Analysis between Features and Target
df_corr = pd.concat([df_features, df_target], axis=1).corr()
df_corr['increase_stock'].sort_values(key=abs, ascending=False)
```

```
[18]: increase_stock      1.000000
hour_of_day_cos     -0.339960
temp                 0.336981
humidity              -0.308726
hour_of_day_sin    -0.308121
summertime            0.216052
month_cos             -0.169059
dew                  0.132663
weekday               -0.116446
visibility            0.113443
windspeed              0.096011
month_sin              -0.092078
day_of_week_sin     -0.088152
precip                -0.059304
snowdepth              -0.047526
cloudcover             -0.045534
day_of_week_cos     -0.031473
holiday                -0.004909
snow                  NaN
Name: increase_stock, dtype: float64
```

Based on above result, The variables hour_of_day_cos, temp, humidity, hour_of_day_sin, and summertime held the biggest correlation to the target variable. Because of this, we could check the trends more based on those attributes

```
[19]: from imblearn.over_sampling import SMOTE

# Applying SMOTE to Balance the Dataset
smote = SMOTE(sampling_strategy='minority')

print("Before SMOTE:\n", df_target.value_counts())
```

```

print()

X_resampled, y_resampled = smote.fit_resample(df_features, df_target)

print("After SMOTE:\n", y_resampled.value_counts())

# Note: Don't use SMOTE outside training data to avoid data leakage, this block
↳ is just for experimentation purposes. In practice, SMOTE should only be
↳ applied to the training set within each cross-validation folds.

```

Before SMOTE:

```

increase_stock
0    1312
1     288
Name: count, dtype: int64

```

After SMOTE:

```

increase_stock
0    1312
1    1312
Name: count, dtype: int64

```

1.2 Model Experiment

```
[20]: from sklearn.metrics import accuracy_score, precision_score, recall_score,
      f1_score, roc_auc_score, classification_report
from sklearn.model_selection import KFold, train_test_split
from sklearn.model_selection import GridSearchCV
```

K-Fold Classification Report

```
[21]: ## Function to do K-Fold Cross Validation and return the classification report
def kfold_classification_reports(clf, X, y, n_splits=10, shuffle=True,
                                 random_state=42):
    kf = KFold(n_splits=n_splits, shuffle=shuffle, random_state=random_state)
    reports = []

    for train_index, test_index in kf.split(X):
        # Applying SMOTE to Balance the Dataset
        smote = SMOTE(sampling_strategy='minority')

        X_train, X_test = X.iloc[train_index], X.iloc[test_index]
        y_train, y_test = y.iloc[train_index], y.iloc[test_index]

        X_resampled, y_resampled = smote.fit_resample(X_train, y_train)
```

```

clf.fit(X_resampled, y_resampled)
y_pred = clf.predict(X_test)

report = classification_report(y_test, y_pred, output_dict=True)
reports += [report]

detailed = []
for i, rep in enumerate(reports):
    fold_result = {
        "fold": i + 1,
        "accuracy": rep["accuracy"],
        "precision": rep["weighted avg"]["precision"],
        "recall": rep["weighted avg"]["recall"],
        "f1": rep["weighted avg"]["f1-score"],
    }
    detailed.append(fold_result)

# compute averaged (generalized) metrics
accuracies = [d["accuracy"] for d in detailed]
precisions = [d["precision"] for d in detailed]
recalls = [d["recall"] for d in detailed]
f1s = [d["f1"] for d in detailed]

report_dict = {
    "general": {
        "accuracy_mean": float(np.mean(accuracies)),
        "accuracy_std": float(np.std(accuracies)),
        "precision_mean": float(np.mean(precisions)),
        "precision_std": float(np.std(precisions)),
        "recall_mean": float(np.mean(recalls)),
        "recall_std": float(np.std(recalls)),
        "f1_mean": float(np.mean(f1s)),
        "f1_std": float(np.std(f1s)),
    },
    "detailed": detailed
}

return report_dict

```

1.3 Benchmark Model

For the benchmark model, we use a naive model that predict each instance as the majority class in the training dataset. This will provide a baseline accuracy to compare the performance of more sophisticated models.

```
[22]: from sklearn.dummy import DummyClassifier
```

```

clf_dummy = DummyClassifier(strategy="stratified")

[23]: reports = kfold_classification_reports(clf_dummy, df_features, df_target)
reports['general']

[23]: {'accuracy_mean': 0.47625,
       'accuracy_std': 0.030593095626301036,
       'precision_mean': 0.6952224635705411,
       'precision_std': 0.032953301949774735,
       'recall_mean': 0.47625,
       'recall_std': 0.030593095626301036,
       'f1_mean': 0.5351422005856692,
       'f1_std': 0.02862709375861207}

```

1.4 Logistic Regression

For Logistic Regression, we will do hyperparameter tuning on Regularization Strength (C), Penalty type (*penalty*), and Solver type (*solver*)

```

[ ]: from sklearn.linear_model import LogisticRegression

clf = LogisticRegression()

#Grid Search to find the best hyperparameters for Logistic Regression
grid_search = GridSearchCV(
    estimator=clf,
    param_grid={
        'C': [0.01, 0.1, 1, 10, 100],
        'penalty': ['l1', 'l2'],
        'solver': ['liblinear', 'lbfgs', 'saga', 'newton-cg']
    },
    scoring='f1',
    cv=10,
    n_jobs=-1,
    verbose=1
)
clf_grid_lg = grid_search.fit(df_features, df_target)

```

```
[249]: clf_grid_lg.best_params_
```

```
[249]: {'C': 10, 'penalty': 'l2', 'solver': 'lbfgs'}
```

```
[250]: clf_lg = LogisticRegression(**clf_grid_lg.best_params_)
```

```
[251]: reports = kfold_classification_reports(clf_lg, df_features, df_target)
reports['general']
```

```
[251]: {'accuracy_mean': 0.8099999999999999,
         'accuracy_std': 0.026545950726994134,
         'precision_mean': 0.8695895287253285,
         'precision_std': 0.016029365215753214,
         'recall_mean': 0.8099999999999999,
         'recall_std': 0.026545950726994134,
         'f1_mean': 0.8271489924908391,
         'f1_std': 0.02043318475431806}
```

1.5 Linear Discriminant Analysis (LDA)

For LDA, we use Hyperparameter tuning for Solver Type (*solver*)

```
[252]: from sklearn.discriminant_analysis import LinearDiscriminantAnalysis

clf = LinearDiscriminantAnalysis()

# Grid Search to find the best hyperparameters for LDA
grid_search = GridSearchCV(
    estimator=clf,
    param_grid={
        'solver': ['svd', 'lsqr', 'eigen']
    },
    scoring='f1',
    cv=10,
    n_jobs=-1,
    verbose=1
)
clf_grid_lda = grid_search.fit(df_features, df_target)
```

Fitting 10 folds for each of 3 candidates, totalling 30 fits

```
[253]: clf_grid_lda.best_params_
```

```
[253]: {'solver': 'svd'}
```

```
[254]: clf_lda = LinearDiscriminantAnalysis(**clf_grid_lda.best_params_)
clf_lda
```

```
[254]: LinearDiscriminantAnalysis()
```

```
[255]: reports = kfold_classification_reports(clf_lda, df_features, df_target)
reports['general']
```

```
[255]: {'accuracy_mean': 0.795625,
         'accuracy_std': 0.0271065605527518,
         'precision_mean': 0.8675792283319778,
         'precision_std': 0.013589091958582378,
```

```
'recall_mean': 0.795625,
'recall_std': 0.0271065605527518,
'f1_mean': 0.8155870764945796,
'f1_std': 0.020417818352304083}
```

1.6 K Nearest Neighbor (KNN)

For K Nearest Neighbor, we use Hyperparameter tuning for selecting the number of neighbors (*n_neighbors*), weight function used in prediction (*weights*), and the metric used for distance computation (*metric*)

```
[256]: from sklearn.neighbors import KNeighborsClassifier

clf = KNeighborsClassifier()

#Grid Search to find the best hyperparameters for KNN
grid_search = GridSearchCV(
    estimator=clf,
    param_grid={
        'n_neighbors': range(3, 21, 2),
        'weights': ['uniform', 'distance'],
        'metric': ['euclidean', 'manhattan', 'minkowski']
    },
    scoring='f1',
    cv=10,
    n_jobs=-1,
    verbose=1
)
clf_grid_knn = grid_search.fit(df_features, df_target)
```

Fitting 10 folds for each of 54 candidates, totalling 540 fits

```
[257]: clf_grid_knn.best_params_
```

```
[257]: {'metric': 'manhattan', 'n_neighbors': 3, 'weights': 'distance'}
```

```
[258]: clf_knn = KNeighborsClassifier(**clf_grid_knn.best_params_)
clf_knn
```

```
[258]: KNeighborsClassifier(metric='manhattan', n_neighbors=3, weights='distance')
```

```
[259]: reports = kfold_classification_reports(clf_knn, df_features, df_target)
reports['general']
```

```
[259]: {'accuracy_mean': 0.7675,
'accuracy_std': 0.020116846174288848,
'precision_mean': 0.823321421215026,
'precision_std': 0.026099143909350715,
```

```
'recall_mean': 0.7675,
'recall_std': 0.020116846174288848,
'f1_mean': 0.7863382939494521,
'f1_std': 0.020420617255737072}
```

1.7 Random Forest Classifier

For Random Forest Classifier, we use Hyperparameter tuning for selecting number of trees (*n_estimator*), maximum depth of each trees (*max_depth*), minimum number of samples required to split (*min_samples_split*), and minimum samples required to be a leaf node (*min_samples_leaf*)

```
[260]: from sklearn.ensemble import RandomForestClassifier
clf = RandomForestClassifier()
grid_search = GridSearchCV(
    estimator=clf,
    param_grid={
        'n_estimators': [20, 50, 100, 200, 250],
        'max_depth': [3, 5, 8, 10, 15, 20],
        'min_samples_split': [2, 5, 10],
        'min_samples_leaf': [1, 2, 4],
        'criterion': ['gini', 'entropy']
    },
    scoring='f1',
    cv=10,
    n_jobs=-1,
    verbose=1
)
clf_grid_rf = grid_search.fit(df_features, df_target)
```

Fitting 10 folds for each of 540 candidates, totalling 5400 fits

```
[261]: clf_grid_rf.best_params_
```

```
[261]: {'criterion': 'entropy',
'max_depth': 15,
'min_samples_leaf': 1,
'min_samples_split': 2,
'n_estimators': 250}
```

```
[262]: clf_rf = RandomForestClassifier(**clf_grid_rf.best_params_)
clf_rf
```

```
[262]: RandomForestClassifier(criterion='entropy', max_depth=15, n_estimators=250)
```

```
[263]: reports = kfold_classification_reports(clf_rf, df_features, df_target)
reports['general']
```

```
[263]: {'accuracy_mean': 0.9037500000000002,
         'accuracy_std': 0.02186606960566987,
         'precision_mean': 0.9065025661469864,
         'precision_std': 0.019154309103939216,
         'recall_mean': 0.9037500000000002,
         'recall_std': 0.02186606960566987,
         'f1_mean': 0.9043197968362898,
         'f1_std': 0.02048728351994524}
```

1.8 Gradient Boosting

For Gradient Boosting, we use Hyperparameter tuning to determine, the number of sequential estimator (*n_estimators*), the learning rate (*learning_rate*), Maximum depth of the individual regression estimators (*max_depth*), and the fraction of samples to be used for fitting the individual estimators (*subsample*).

```
[264]: from sklearn.ensemble import GradientBoostingClassifier
clf = GradientBoostingClassifier()
grid_search = GridSearchCV(
    estimator=clf,
    param_grid={
        'n_estimators': [50, 100, 200],
        'learning_rate': [0.01, 0.05, 0.1, 0.2, 0.5],
        'max_depth': [3, 5, 8, 10],
        'subsample': [0.6, 0.8, 1.0]
    },
    scoring='f1',
    cv=10,
    n_jobs=-1,
    verbose=1
)
clf_grid_gb = grid_search.fit(df_features, df_target)
```

Fitting 10 folds for each of 180 candidates, totalling 1800 fits

```
[267]: clf_grid_gb.best_params_
```

```
[267]: {'learning_rate': 0.1, 'max_depth': 10, 'n_estimators': 100, 'subsample': 0.6}
```

```
[268]: clf_gb = GradientBoostingClassifier(**clf_grid_gb.best_params_)
clf_gb
```

```
[268]: GradientBoostingClassifier(max_depth=10, subsample=0.6)
```

```
[269]: reports = kfold_classification_reports(clf_gb, df_features, df_target)
reports['general']
```

```
[269]: {'accuracy_mean': 0.899375,
'accuracy_std': 0.02043472840533977,
'precision_mean': 0.9014960071645403,
'precision_std': 0.018078538997051353,
'recall_mean': 0.899375,
'recall_std': 0.02043472840533977,
'f1_mean': 0.8998663647708327,
'f1_std': 0.01972180961714667}
```

1.8.1 Saving Model

Saving all model so we can use it later

```
[270]: # Applying SMOTE to Balance the Full Training Dataset
smote = SMOTE(sampling_strategy='minority')

print("Before SMOTE:\n", df_target.value_counts())

print()

X_resampled, y_resampled = smote.fit_resample(df_features, df_target)

print("After SMOTE:\n", y_resampled.value_counts())
```

```
Before SMOTE:
increase_stock
0    1312
1     288
Name: count, dtype: int64

After SMOTE:
increase_stock
0    1312
1    1312
Name: count, dtype: int64
```

```
[271]: import pickle

# Save the Dummy Classifier model used as benchmark
with open('clf_dummy.pkl', 'wb') as f:
    pickle.dump(clf_dummy, f)
```

```
[272]: # Train the Logistic Regression model on the full training dataset
clf_lg = LogisticRegression(**clf_grid_lg.best_params_)
clf_lg.fit(X_resampled, y_resampled)

# Save the trained Logistic Regression model
with open('clf_lg.pkl', 'wb') as f:
```

```

pickle.dump(clf_lg, f)

[ ]: # Train the LDA model on the full training dataset
clf_lda = LinearDiscriminantAnalysis(**clf_grid_lda.best_params_)
clf_lda.fit(X_resampled, y_resampled)

# Save the trained LDA model
with open('clf_lda.pkl', 'wb') as f:
    pickle.dump(clf_lda, f)

[274]: # Train the KNN model on the full training dataset
clf_knn = KNeighborsClassifier(**clf_grid_knn.best_params_)
clf_knn.fit(X_resampled, y_resampled)

# Save the trained KNN model
with open('clf_knn.pkl', 'wb') as f:
    pickle.dump(clf_knn, f)

[275]: # Train the Random Forest Classifier model on the full training dataset
clf_rf = RandomForestClassifier(**clf_grid_rf.best_params_)
clf_rf.fit(X_resampled, y_resampled)

# Save the trained Random Forest Classifier model
with open('clf_rf.pkl', 'wb') as f:
    pickle.dump(clf_rf, f)

[276]: # Train the Gradient Boosting Classifier model on the full training dataset
clf_gb = GradientBoostingClassifier(**clf_grid_gb.best_params_)
clf_gb.fit(X_resampled, y_resampled)

# Save the trained Gradient Boosting Classifier model
with open('clf_gb.pkl', 'wb') as f:
    pickle.dump(clf_gb, f)

[ ]:

```